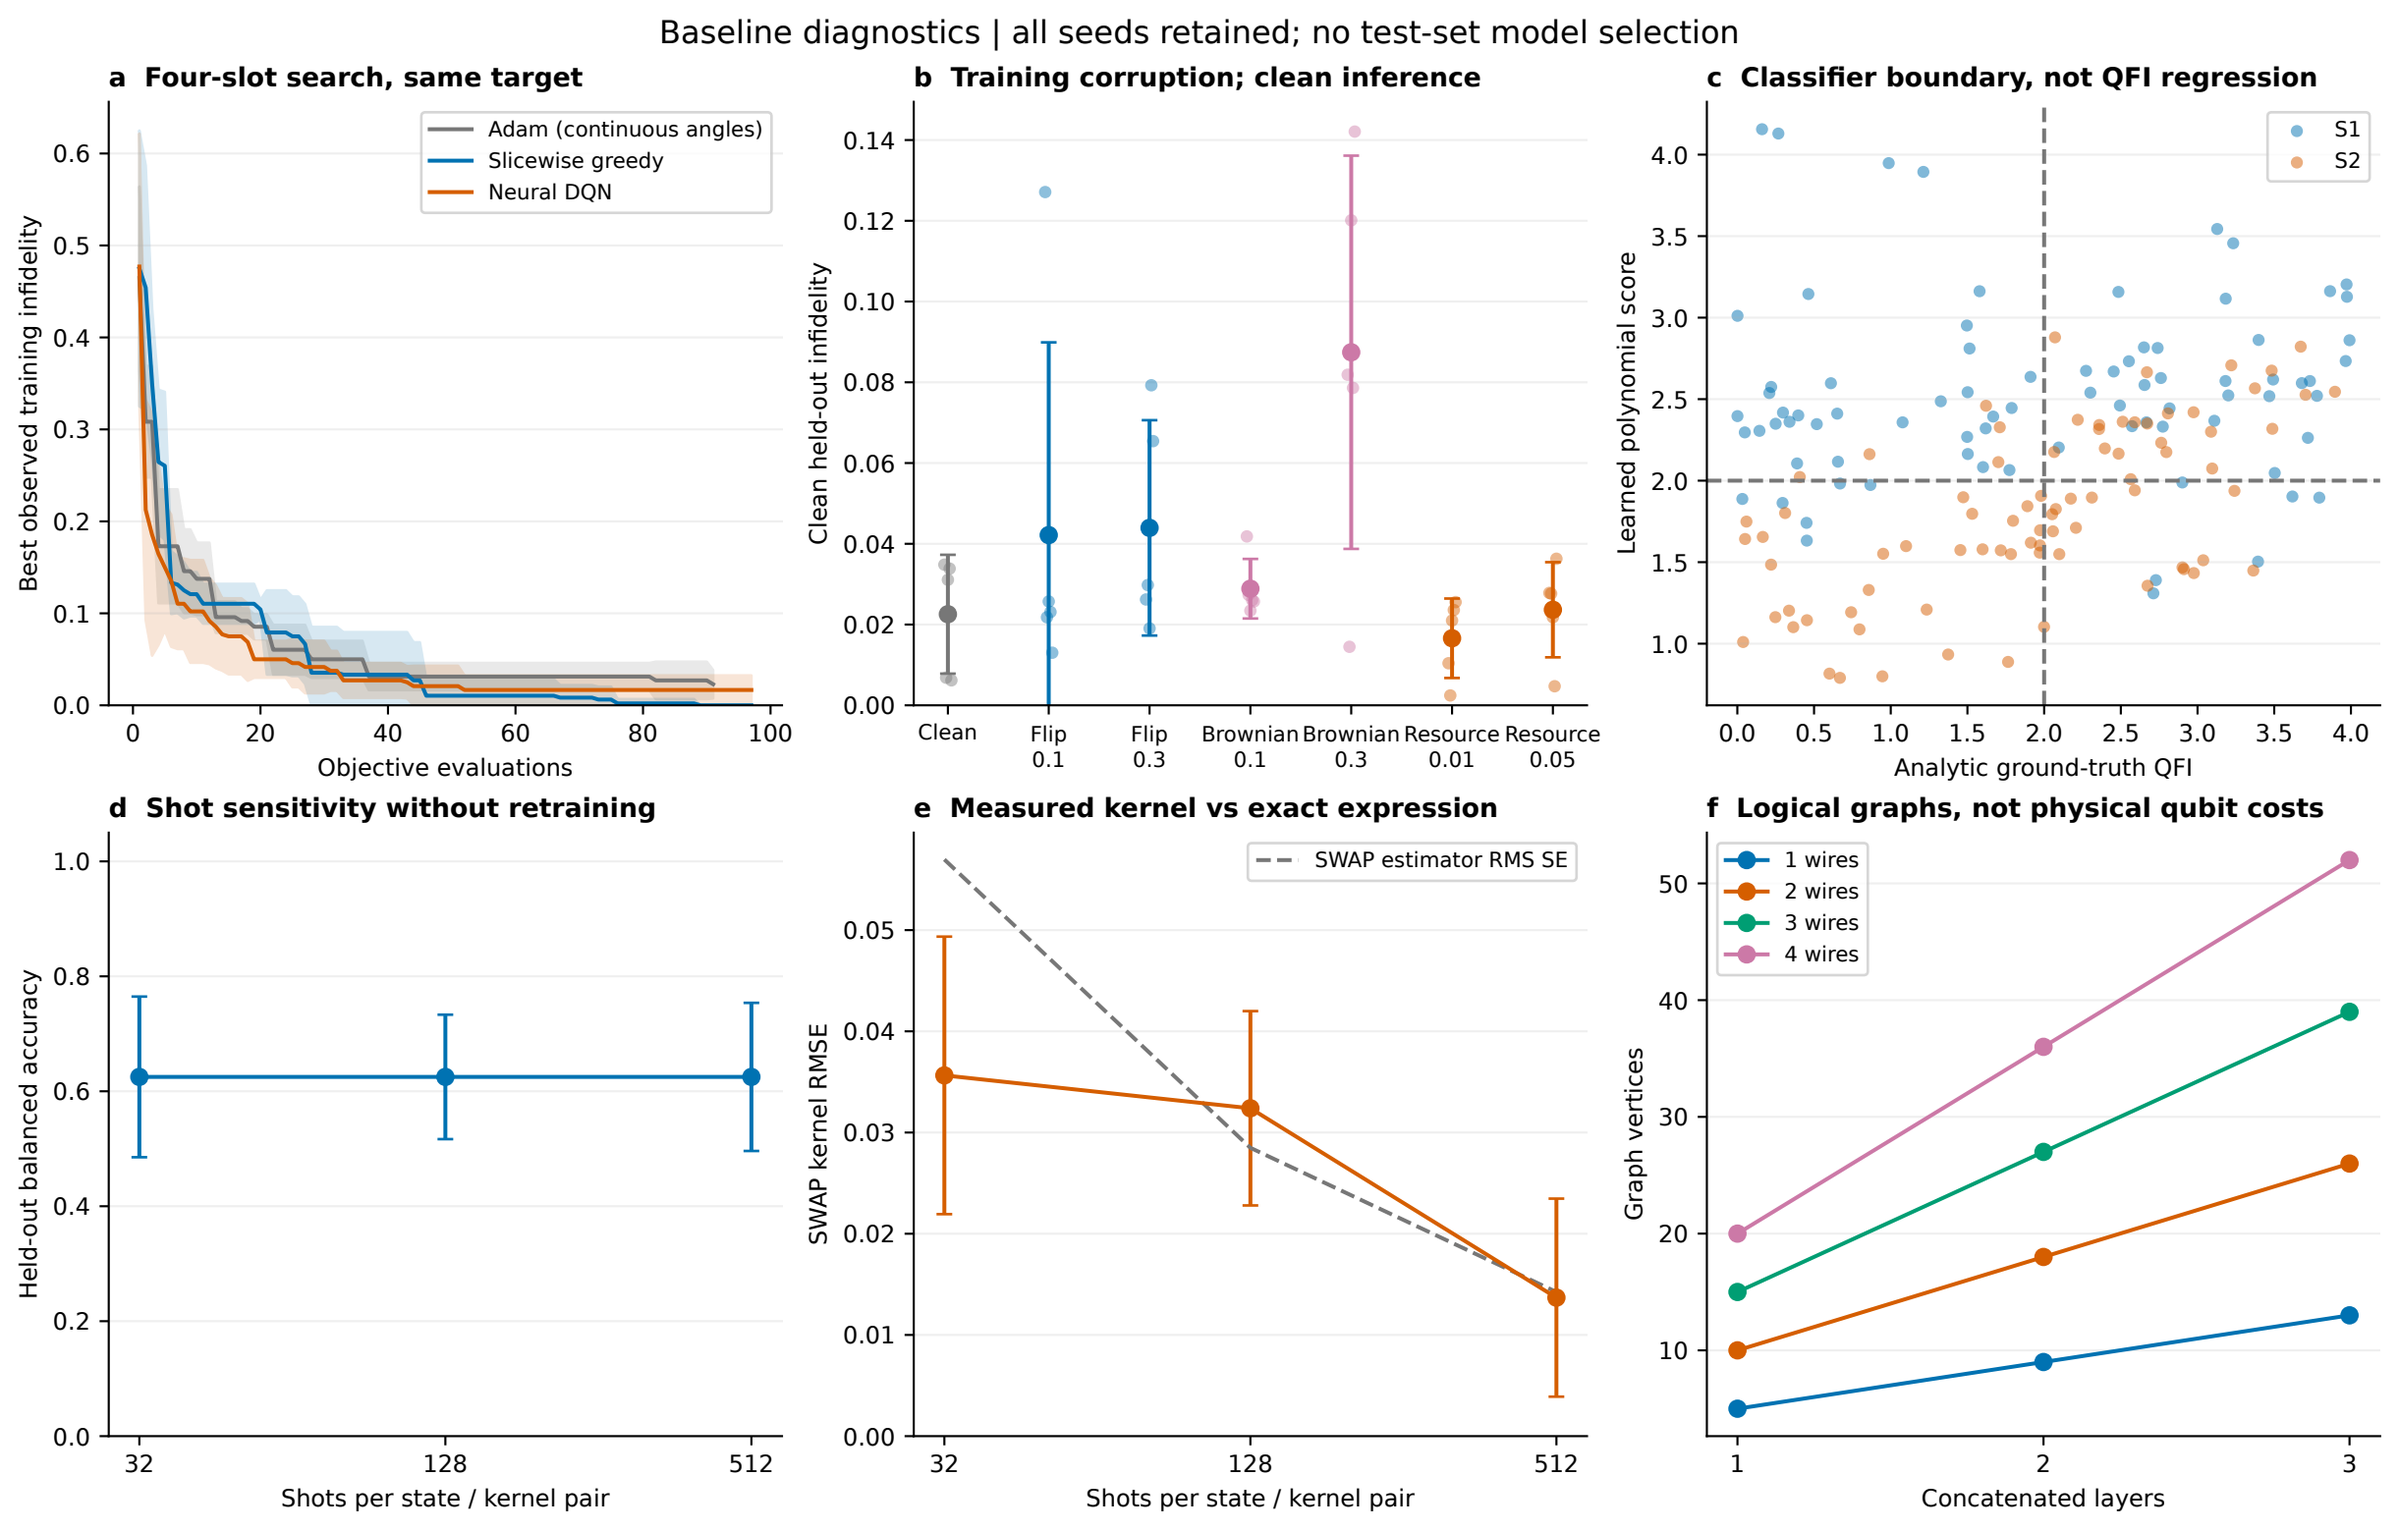
Baseline diagnostics | all seeds retained; no test-set model selection
a Four-slot search, same target
Adam (continuous angles)
Slicewise greedy
Neural DQN
Best observed training infidelity
Objective evaluations
b Training corruption; clean inference
Clean held-out infidelity
Clean
Flip 0.1
Flip 0.3
Brownian 0.1
Brownian 0.3
Resource 0.01
Resource 0.05
c Classifier boundary, not QFI regression
S1
S2
Learned polynomial score
Analytic ground-truth QFI
d Shot sensitivity without retraining
Held-out balanced accuracy
Shots per state / kernel pair
e Measured kernel vs exact expression
SWAP estimator RMS SE
SWAP kernel RMSE
Shots per state / kernel pair
f Logical graphs, not physical qubit costs
1 wires
2 wires
3 wires
4 wires
Graph vertices
Concatenated layers